

THE GPU SHOWDOWN

An evidence dossier for curious computer scientists

Plot Twist. Something is wrong with our GPU ranking. The newest card did not always win. One assignment was stopped before it began because the computer was not the GPU we thought it was. Your mission is to decide what the evidence actually supports.

The case file

We measured Aider/Ollama inference on a small fleet. The headline number is *observed whole-system inference latency*: wall-clock time shaped by the GPU, memory, model residency, CPU, operating system, backend, driver, load, and experimental conditions. It is not a silicon-only score.

The useful question is not “Which GPU wins?” It is: *What claim can this evidence carry, and what better test would make the claim stronger?*

Evidence Check. The automated percentage in the source reports is a concept coverage screen made from keyword/pattern checks. It is not a validated correctness grade.

Meet the machines

Machine	GPU	Family	VRAM (GB)	Warm s	Status
April	RTX 5080	Blackwell	16	2.48	safe
Maise	RTX 2080 Super	Turing	8	5.79	safe
Faith	GTX 1660 Ti Max-Q	Turing	6	7.02	safe
Brandy	Tesla T4	Turing	15	9.69	safe
STF125	RTX 5060 Ti	Blackwell	8	10.80	caveated
STF259	RTX 3060 Ti	Ampere	8	12.51	caveated
Grace	GTX 1060 3GB	Pascal	3	14.73	caveated
Morgan	GTX 1080	Pascal	8	15.35	caveated

The values are curated from source experiment directories or immutable source-import packets, not copied from stale aggregate HTML. STF125 has 84 successful calls and 6 explicit timeouts; STF259 has 90 successful calls and no errors. Grace has a residency limitation; Morgan’s context includes WSL2. All cross-machine comparisons remain caveated.

Confound Ahead. Maise’s Turing RTX 2080 Super appears faster than the newer Blackwell RTX 5060 Ti and Ampere RTX 3060 Ti in these observed whole-system runs. That is a question generator, not proof that Turing is inherently faster than either newer architecture.

How the experiment worked

The original showdown used three prompt levels, five repetitions, six local models, wall-clock timing, raw outputs, hardware provenance, and an automated concept screen. A complete run targeted 90 calls. A timeout is a state, not a zero; a clean exit is not proof that an explanation is correct.

Supported observation	Claim we cannot make
Recorded systems had different observed latencies	A pure GPU-generation multiplier
Grace’s setup has a 3 GB residency limitation	An exact offload penalty without a controlled test
The gate found the wrong hardware	A numeric RTX 5060 benchmark

Trust Nothing. A direct hardware check stopped the intended “STF125 RTX 5060” assignment when the machine was actually an RTX 4070 with about 12 GB VRAM. Refusing to run the wrong experiment is a successful scientific outcome.

DSCT lens: from observation to warrant

Use the chain **OBSERVATION** → **CLAIM** → **EVIDENCE** → **WARRANT** → **CONFOUND** → **BETTER TEST**.

1. Write one observation from the table.
2. Turn it into a modest claim.
3. Cite the source row/path.
4. State the warrant linking evidence to claim.
5. Name a confound.
6. Design a better test.

The wrong-hardware gate and timeout counts are evidence about experimental discipline, not missing zeros. Explain why “newer always wins” is tempting but unsupported here.